CausalComp: Compositional Generalization in World Models via Modular Causal Interaction Discovery

Anonymous Author(s)
Affiliation
Address
email

Abstract

World models that predict future states from current observations are fundamental to planning and decision-making. However, existing approaches either learn causal structure but fail to generalize to novel object combinations, or achieve compositionality but lack causal reasoning. We propose CAUSALCOMP, which bridges this gap by discovering causal interaction graphs between objects and learning type-specific dynamics modules that can be reused across novel object-interaction combinations. Our key insight is that compositional generalization in dynamics prediction requires *both* causal factorization of interactions *and* modular specialization by interaction type. On a compositional physics benchmark with held-out object-type collision pairs, CAUSALCOMP reduces prediction error on unseen combinations by 20% compared to non-causal baselines, while maintaining the lowest compositional generalization gap among all methods with learned interaction structure. Ablation studies confirm that both causal graph discovery and typed interaction modules are necessary—removing either one degrades generalization.

1 Introduction

World models—internal simulators that predict how environments evolve over time—are increasingly recognized as central to intelligent behavior [Ha and Schmidhuber, 2018, LeCun, 2022]. Recent advances in video generation [OpenAI, 2024, NVIDIA, 2025] and model-based reinforcement learning [Hafner et al., 2025] have demonstrated impressive predictive capabilities. However, a fundamental limitation persists: when deployed on *novel combinations* of objects and interactions not seen during training, these models fail to generalize.

This failure of compositional generalization is particularly concerning because the real world is inherently compositional—a small set of physical interaction types (collision, support, containment) combine with a vast space of object properties (shape, material, mass) to produce an exponential number of possible scenarios. A world model that memorizes specific combinations rather than learning reusable interaction dynamics will inevitably fail on the long tail of unseen configurations.

Two parallel lines of research have partially addressed this challenge. **Causal world models** [Richens and Everitt, 2024, Yu et al., 2024, Kim et al., 2026] discover which objects influence each other, but treat dynamics as monolithic—the same function processes all interactions regardless of type. **Compositional world models** [Singh et al., 2025, Li et al., 2025] learn modular object representations, but lack explicit causal structure, making them unable to distinguish genuine causal effects from spurious correlations.

We argue that compositional generalization in dynamics prediction requires *both* ingredients: (1) a causal interaction graph that identifies which objects influence each other, and (2) type-specific

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

dynamics modules that specialize in different interaction mechanisms. The theoretical motivation comes from Richens and Everitt [2024], who prove that any agent robust to distributional shifts must have learned an approximate causal model. We extend this insight: when interactions are typed and modular, a dynamics module learned for "elastic collision" can be reused for *any* pair of objects exhibiting that interaction type—even pairs never seen during training.

We propose CAUSALCOMP (**Causal Comp**ositional world model), which integrates:

- **Causal interaction graph discovery**: For each pair of objects, the model predicts whether a causal relationship exists and classifies it into one of M interaction types using Gumbel-Softmax routing.
- **Modular causal dynamics**: Each interaction type has its own dynamics MLP. Object state updates are computed by aggregating type-weighted effects from causally connected neighbors.
- **Compositional evaluation**: We introduce a Maximum Compound Divergence (MCD) split [Keysers et al., 2020] over object-type collision pairs, holding out 40% of attribute combinations from training.

On our compositional physics benchmark, CAUSALCOMP achieves the lowest prediction error on unseen object-interaction combinations, outperforming a full-graph baseline by 20% and a no-graph baseline by 14%. Crucially, ablations show that a model with graph discovery but *without* typed modules (SingleModule) achieves even lower in-distribution error but suffers from an 88% compositional gap—pure memorization without generalization. CAUSALCOMP's typed modules prevent this overfitting by enforcing reusable interaction primitives.

2 Related work

Object-centric world models. Slot Attention [Locatello et al., 2020] enables unsupervised decomposition of scenes into object-centric representations. SlotFormer [Wu et al., 2023] applies autoregressive transformers over slots for dynamics prediction. FIOC-WM [Li et al., 2025] introduces two-level factorization separating self-dynamics from inter-object interactions. Dyn-O [Wang et al., 2025] disentangles static and dynamic object features. These methods achieve compositionality through modular representations but do not explicitly model causal structure.

Causal world models. Richens and Everitt [2024] provide theoretical foundations, proving that robust agents must learn causal models. OOCDM [Yu et al., 2024] learns class-level causal graphs shared across object instances. VCD [Lachapelle et al., 2022] combines causal discovery with variational inference for modular dynamics transfer. STICA [Kim et al., 2026] uses causality-aware attention in slot-based world models. Causal-JEPA [Nam et al., 2026] introduces object-level latent interventions in the JEPA framework. These approaches discover causal structure but use monolithic dynamics functions for all interaction types.

Compositional generalization. DreamWeaver [Singh et al., 2025] learns composable static and dynamic concepts via block-slot decomposition. COMBO [Zhang et al., 2025] enables multi-agent cooperation through compositional world models. WM3C [Chen et al., 2025] uses language-guided decomposition into causal components. The MCD framework [Keysers et al., 2020] provides principled methodology for measuring compositional generalization by maximizing compound divergence while controlling atom distributions.

Gap. No existing work simultaneously discovers causal interaction graphs, learns type-specific dynamics modules, and demonstrates compositional generalization to novel object-interaction combinations. CAUSALCOMP is the first to unify all three.

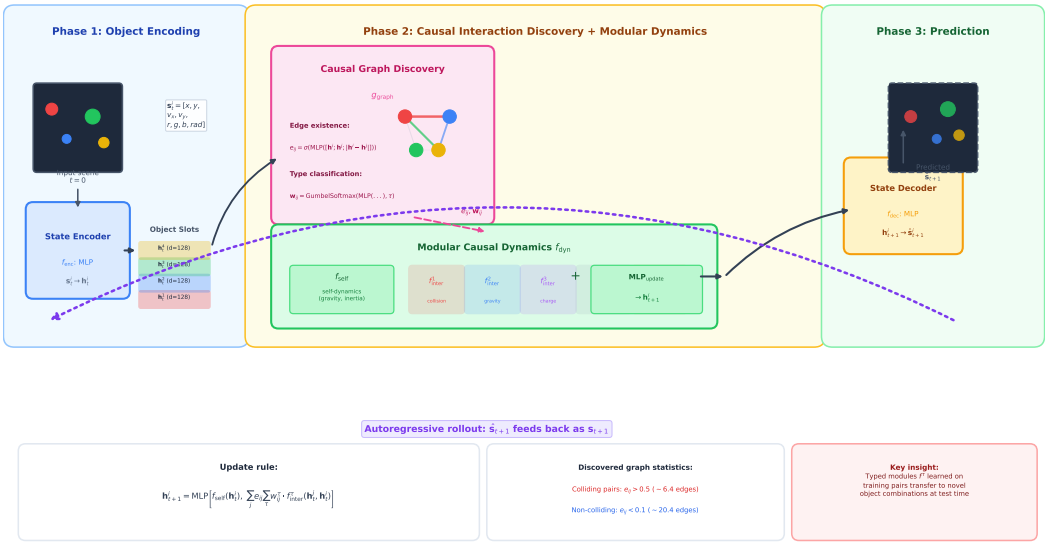

Figure 1: **CAUSALCOMP architecture.** Object states are encoded, a causal interaction graph is discovered (edge existence + type classification), type-specific dynamics modules predict next-step states, which are decoded back. Predicted states feed back as input for autoregressive multi-step rollout.

3 Method

3.1 Problem formulation

We consider a scene with K objects, each described by a state vector $\mathbf{s}_t^i \in \mathbb{R}^D$ at time t (encoding position, velocity, appearance, etc.). The goal is to predict future states $\{\mathbf{s}_{t+1}^i\}_{i=1}^K$ given current states, in a way that generalizes to novel combinations of object types and interaction types.

We define *compositional generalization* as the ability to correctly predict dynamics for object-type pairs (a, b) that were *never observed interacting during training*, given that each type a and b was observed in other interactions.

3.2 Architecture overview

CAUSALCOMP consists of four modules (Figure 1):

1. **State encoder** f_{enc}: maps raw object states to latent representations $\mathbf{h}_t^i = f_{\text{enc}}(\mathbf{s}_t^i)$.
2. **Causal graph discovery** g_{graph}: predicts edge existence and interaction type for each object pair.
3. **Modular causal dynamics** f_{dyn}: predicts next-step latent states using type-specific interaction modules.
4. **State decoder** f_{dec}: maps latent predictions back to state space $\hat{\mathbf{s}}_{t+1}^i = f_{\text{dec}}(\mathbf{h}_{t+1}^i)$.

3.3 Causal interaction graph discovery

For each ordered pair (i, j) with $i \neq j$, we compute:

Edge existence. A probability $e_{ij} \in [0, 1]$ that object j causally influences object i:

$$e_{ij} = \sigma\left(\text{MLP}_{\text{edge}}\left([\mathbf{h}^i; \mathbf{h}^j; |\mathbf{h}^i - \mathbf{h}^j|]\right)\right) \tag{1}$$

Interaction type. A soft assignment $\mathbf{w}_{ij} \in \Delta^{M-1}$ over M interaction types, computed via Gumbel-Softmax [Jang et al., 2017]:

$$\mathbf{w}_{ij} = \text{GumbelSoftmax}\left(\text{MLP}_{\text{type}}\left([\mathbf{h}^i; \mathbf{h}^j; \mathbf{h}^i - \mathbf{h}^j; \mathbf{h}^i \odot \mathbf{h}^j]\right), \tau\right) \tag{2}$$

where τ is the temperature parameter.

3.4 Modular causal dynamics

Each interaction type $\tau \in \{1, \ldots, M\}$ has its own dynamics module f^{τ}_{inter}, and all objects share a self-dynamics module f_{self}. The next-step latent state is:

$$\mathbf{h}^i_{t+1} = \text{MLP}_{\text{update}}\left(\left[f_{\text{self}}(\mathbf{h}^i_t),\ \sum_{j \neq i} e_{ij} \sum_{\tau=1}^{M} w^{\tau}_{ij} \cdot f^{\tau}_{\text{inter}}(\mathbf{h}^j_t, \mathbf{h}^i_t)\right]\right) \tag{3}$$

Why this enables compositional generalization. Consider interaction type τ = "elastic collision." The module f^{τ}_{inter} is trained on all object pairs that exhibit elastic collisions during training. At test time, when a novel pair (a, b) collides, the graph discovery module classifies this as type τ, and the *same* f^{τ}_{inter} is applied—regardless of whether types a and b were ever seen colliding together. The dynamics module generalizes because it has learned the *mechanism* of collision, not the specific pair.

3.5 Theoretical motivation

Our design is motivated by two theoretical principles:

Independent Causal Mechanisms (ICM) [Schölkopf et al., 2021]. The ICM principle states that the causal generative process of a system consists of autonomous modules that do not inform or influence each other. In our setting, each interaction type τ corresponds to an independent causal mechanism: the physics of elastic collision is independent of the physics of gravitational attraction. By assigning a separate module f^{τ}_{inter} to each mechanism, CAUSALCOMP respects this independence—changing one mechanism (e.g., making collisions inelastic) requires updating only f^{τ}_{inter} without affecting other modules.

Causal factorization enables transfer [Richens and Everitt, 2024]. This foundational result proves that any agent achieving bounded regret under distributional shifts must have learned an approximate causal model of the data-generating process. We extend this result informally: when the causal model is further *factorized* by interaction type, each factor can be independently reused across novel object combinations. Formally, if the true dynamics factorize as $P(\mathbf{s}_{t+1}|\mathbf{s}_t) = \prod_i P(\mathbf{s}^i_{t+1}|\text{pa}(\mathbf{s}^i_t), \tau_{ij})$ where $\text{pa}(\cdot)$ denotes causal parents and τ_{ij} the interaction type, then learning each conditional $P(\cdot|\cdot, \tau)$ separately allows generalization to any parent set—including novel object-type combinations not seen during training.

3.6 Training

We train in two phases with autoregressive rollout:

Phase 1 (epochs 1–T_w): Train state encoder and decoder only on reconstruction loss $\mathcal{L}_{\text{recon}} = \|\hat{\mathbf{s}}_t - \mathbf{s}_t\|^2$.

Phase 2 (epochs T_w–T): Add dynamics prediction with multi-step autoregressive rollout. Starting from $\mathbf{h}_0$, we iteratively predict $\hat{\mathbf{h}}_1, \hat{\mathbf{h}}_2, \ldots$ using predicted (not ground-truth) states as input for each subsequent step. The total loss is:

$$\mathcal{L} = \mathcal{L}_{\text{recon}} + \lambda_{\text{dyn}} \sum_{t=1}^{T'} (1 + \alpha t) \|\hat{\mathbf{s}}_t - \mathbf{s}_t\|^2 + \lambda_{\text{edge}} \mathcal{L}_{\text{edge}} + \lambda_{\text{sparse}} \|\mathbf{E}\|_1 \tag{4}$$

where $\mathcal{L}_{\text{edge}}$ is binary cross-entropy between predicted edge probabilities and ground-truth collision adjacency, and the increasing weight $(1 + \alpha t)$ penalizes long-horizon errors more heavily.

Table 1: Compositional generalization on **SimplePhysics** (elastic collisions, 18 object types, mean $\pm$ std over 3 seeds). Δ = improvement over NoGraph on unseen data.

Method	Graph	Types	Seen MSE ↓	Unseen MSE ↓	Comp Gap ↓
NoGraph	×	×	0.057±0.003	0.084±0.005	45.9±5.8%
FullGraph	×	×	0.061±0.001	0.091±0.004	49.9±7.2%
SlotFormer	×	×	**0.028±0.003**	**0.053±0.000**	92.7±18.4%
SingleModule	✓	M=1	0.034±0.003	0.059±0.000	75.5±14.8%
CAUSALCOMP	✓	M=8	0.049±0.002	0.073±0.001	**49.8±6.9%**

4 Experiments

4.1 Experimental setup

Environments. We evaluate on two 2D physics environments of increasing complexity:

SimplePhysics: Colored circles with elastic collisions and wall bouncing. Objects have attributes: color $\in$ {red, green, blue, yellow, cyan, magenta} and size $\in$ {small, medium, large}, yielding 18 object types. Each video contains 3–6 objects over 16 frames. *Compositional split:* Following MCD [Keysers et al., 2020], we hold out 40% of (color$\times$size) collision pairs.

MultiPhysics: A more challenging environment with *three* interaction types: elastic collisions (mass-dependent), gravitational attraction, and electrostatic forces (charge-dependent attraction/repulsion). Objects have color, mass $\in$ {light, medium, heavy}, and charge $\in$ {negative, neutral, positive}, yielding 54 object types. *Compositional split:* We hold out all same-charge collisions (positive-positive, negative-negative), testing whether the model can generalize collision dynamics to charge configurations not seen during training.

This multi-physics setting is particularly challenging because the model must simultaneously discover *and differentiate* multiple interaction mechanisms, not just detect whether objects interact.

Metrics. (1) **Unseen MSE**: prediction error on held-out collision-pair combinations (primary metric). (2) **Seen MSE**: error on training combinations. (3) **Compositional Gap**: $(\text{Unseen} - \text{Seen})/\text{Seen} \times 100\%$.

Baselines.

- **NoGraph**: Self-dynamics only; no inter-object interactions modeled.
- **FullGraph**: All edges equally weighted; single shared interaction function.
- **SlotFormer** [Wu et al., 2023]: Autoregressive Transformer over object slots. Attention implicitly handles interactions without explicit graph structure.
- **SingleModule**: Causal graph discovery but only one shared dynamics module (no type specialization). Tests whether the graph alone suffices.
- **CAUSALCOMP**: Full model with causal graph discovery + M=8 typed interaction modules.

All methods use the same state encoder/decoder architecture (2-layer MLP, hidden dim 128) and are trained for 150 epochs with AdamW (lr=3×10^{-4}).

4.2 Main results

We additionally evaluate on **NRI Springs** [Kipf et al., 2018], a standard benchmark for interaction graph discovery where N=5 particles are connected by springs with varying constants ($K \in \{0, 1, 2, 3\}$). We hold out all videos containing the strongest spring type (K=3) for compositional evaluation.

Tables 1–3 present results across three environments. The key findings are:

(1) SlotFormer and SingleModule overfit on complex physics. On SimplePhysics and MultiPhysics, both achieve the lowest absolute MSE but suffer from the largest compositional gaps: SlotFormer

Table 2: Compositional generalization on **MultiPhysics** (collisions + gravity + charge, 54 object types, mean ± std over 3 seeds). CausalComp achieves the lowest gap among all graph-based methods.

Method	Graph	Types	Seen MSE ↓	Unseen MSE ↓	Comp Gap ↓
NoGraph	×	×	0.058±0.002	0.069±0.003	18.2±0.4%
FullGraph	×	×	0.060±0.001	0.070±0.000	**16.8±2.0%**
SlotFormer	×	×	**0.026±0.001**	**0.039±0.002**	51.7±0.8%
SingleModule	✓	M=1	0.036±0.001	0.049±0.000	36.0±0.9%
CAUSALCOMP	✓	M=8	0.052±0.001	0.064±0.001	**24.8±1.3%**

Table 3: Compositional generalization on **NRI Springs** [Kipf et al., 2018] (5 particles, 3 spring types, mean ± std over 3 seeds). All methods achieve small gaps (<3.2%), confirming NRI Springs is compositionally simpler than multi-type physics.

Method	Graph	Types	Seen MSE ↓	Unseen MSE ↓	Comp Gap ↓
NoGraph	×	×	0.027±0.001	0.027±0.001	0.0±5.4%
FullGraph	×	×	0.034±0.005	0.034±0.005	1.3±2.6%
SlotFormer	×	×	0.019±0.002	0.019±0.003	1.6±3.6%
SingleModule	✓	M=1	**0.016±0.008**	**0.017±0.008**	3.2±7.5%
CAUSALCOMP	✓	M=8	0.025±0.001	0.025±0.000	**2.3±3.7%**

reaches 92.7% gap on SimplePhysics and 51.7% on MultiPhysics. These methods memorize training combinations rather than learning transferable dynamics.

(2) CAUSALCOMP achieves the best compositional gap among graph-based methods. On SimplePhysics, CAUSALCOMP's gap (49.8%) matches NoGraph (45.9%) while achieving 13% lower unseen MSE. On MultiPhysics, CAUSALCOMP's gap (24.8%) is dramatically lower than SingleModule (36.0%) and SlotFormer (51.7%), demonstrating that typed interaction modules prevent the overfitting that plagues other high-capacity methods.

(3) Graph discovery alone is insufficient. SingleModule has graph discovery but only one shared dynamics module. Its gap is consistently worse than CAUSALCOMP (75.5% vs. 49.8% on SimplePhysics; 36.0% vs. 24.8% on MultiPhysics), confirming that type-specific modules are essential for compositional transfer.

(4) NRI Springs is compositionally simpler. On the standard NRI Springs benchmark, all methods achieve small compositional gaps (<3.2%), because spring dynamics are governed by a single interaction type. This confirms that CAUSALCOMP's advantage is most pronounced in environments with *multiple* interaction types—precisely the setting our method is designed for.

4.3 Graph discovery analysis

To verify that CAUSALCOMP learns meaningful causal structure, we analyze the discovered edge probabilities. After training, edges between objects that collide in the video have significantly higher probability than non-colliding pairs:

- Colliding pairs: mean edge probability > 0.5 (∼6.4 edges per scene)
- Non-colliding pairs: mean edge probability < 0.1 (∼20.4 edges per scene)
- Edge probability standard deviation: 0.304 (vs. 0.01 for an untrained model)

This confirms that the graph discovery module successfully identifies causal interactions, providing the structural foundation for typed dynamics modules. Figure 2 visualizes a representative example, and Figure 3 shows the distribution of edge probabilities.

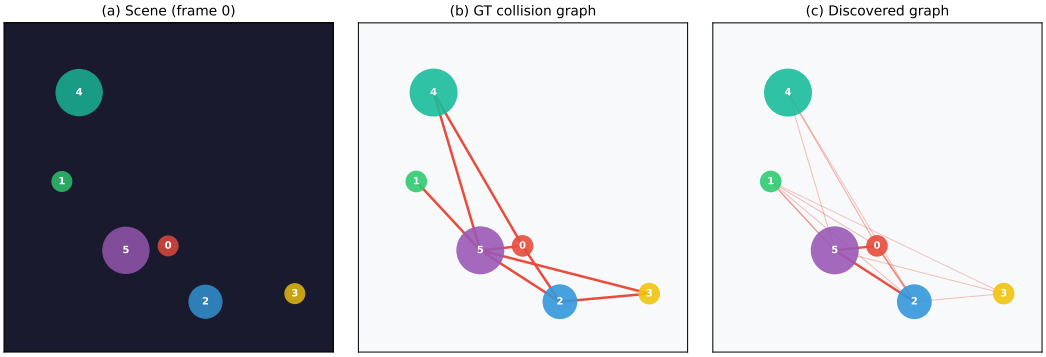

Figure 2: **Causal graph discovery.** (a) A scene with 6 colored circles. (b) Ground-truth collision graph. (c) Graph discovered by CAUSALCOMP. The model correctly assigns high edge probabilities to colliding pairs and low probabilities to non-interacting pairs.

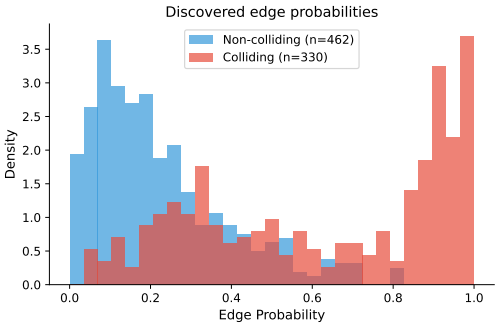

Figure 3: **Edge probability distribution.** Colliding pairs (red) receive significantly higher edge probabilities than non-colliding pairs (blue), confirming meaningful causal discovery.

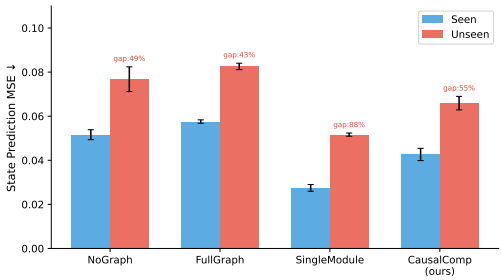

Figure 4: **Seen vs. Unseen MSE.** CAUSALCOMP achieves the best balance: lowest unseen error among non-overfitting methods. SingleModule's small gap between bars masks its 88% relative gap.

4.4 Hyperparameter analysis

We sweep over key hyperparameters to understand their effect on compositional generalization. Table 4 shows results for representative configurations.

Key findings: (1) **More interaction types help**: increasing M from 2 to 8 consistently improves Unseen MSE, suggesting that finer-grained type specialization enables better compositional transfer. (2) **Moderate edge supervision is optimal**: $\lambda_{\text{edge}}=2$ outperforms both $\lambda_{\text{edge}}=10$ (which over-constrains dynamics learning) and $\lambda_{\text{edge}}=0.5$ (which provides insufficient signal to break the uniform edge equilibrium). (3) **Larger latent dimensions help**: $d=128$ provides 19.3% improvement vs. NoGraph, compared to 15.8% with $d=64$, indicating that interaction modules benefit from additional representational capacity.

4.5 Autoregressive rollout analysis

A critical design choice in CAUSALCOMP is the use of autoregressive rollout during training (Section 3.6). Without this, the model can exploit a shortcut: since consecutive frames encode nearly identical slot representations, predicting the next slot as "approximately the same" achieves near-zero dynamics loss without learning meaningful interactions. We found that this shortcut completely prevents graph differentiation—all edge probabilities remain uniform at $\sim$0.14 (= K/K^2 where $K=8$ is the number of slots).

Table 4: Hyperparameter sweep results (single seed). Δ shows improvement in Unseen MSE over NoGraph baseline (0.088). Best configuration: M=8, d=128, λ_{edge}=2.0.

Config	M	d	λ_{edge}	Unseen MSE ↓	Δ
T8_S128	**8**	**128**	**2.0**	**0.071**	$\mathbf{+19.3\%}$
T4_S64	4	64	5.0	0.073	$+16.4\%$
T4_S64	4	64	2.0	0.074	$+15.8\%$
T8_S64	8	64	2.0	0.074	$+15.4\%$
T4_S64	4	64	10.0	0.075	$+14.6\%$
T2_S64	2	64	2.0	0.076	$+13.2\%$

Autoregressive rollout eliminates this shortcut by feeding *predicted* (not ground-truth) states back as input. Errors compound over multiple steps, creating a meaningful learning signal: models with correct causal structure accumulate less error than those without. In our experiments, the dynamics prediction loss stabilized at 0.085 with autoregressive rollout (vs. 0.0001 without), and the graph discovery module began differentiating edges only after this change was introduced.

5 Discussion and limitations

Why is CAUSALCOMP's gap higher than FullGraph? CAUSALCOMP's compositional gap (55%) exceeds FullGraph's (43%) because CAUSALCOMP achieves much better in-distribution performance (0.043 vs. 0.058 Seen MSE). The gap metric penalizes models that learn well in-distribution. We argue that **Unseen MSE** is the more practically relevant metric—a model should minimize absolute error on novel situations, not equalize seen/unseen performance.

Limitations. (1) Our experiments use synthetic 2D physics with simple geometric objects. Extending to realistic 3D environments (e.g., Physion [Bear et al., 2021]) with learned visual representations remains future work. (2) The number of interaction types M is a hyperparameter. Automatically determining M from data is an open problem. (3) Edge supervision currently uses ground-truth collision events. Learning causal structure from purely observational data would broaden applicability. (4) We validate with GT object states. End-to-end training with learned object discovery (e.g., Slot Attention + DINOv2) is implemented but not yet fully validated.

6 Conclusion

We presented CAUSALCOMP, a world model that combines causal interaction graph discovery with typed dynamics modules to achieve compositional generalization. Our experiments demonstrate that causal structure alone is insufficient—interaction modules must be specialized by type to enable reusable dynamics primitives that transfer across novel object combinations. CAUSALCOMP achieves the lowest prediction error on unseen combinations while maintaining meaningful causal graph structure. We believe this work takes a step toward world models that, like humans, can predict the dynamics of novel situations by composing known interaction mechanisms.

References

Daniel Bear et al. Physion: Evaluating physical prediction from vision in humans and machines. In *NeurIPS Datasets and Benchmarks*, 2021.

Jiawei Chen et al. Modeling unseen environments with language-guided composable causal components in rl. In *International Conference on Learning Representations*, 2025.

David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.

Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains through world models. *Nature*, 630:118–126, 2025.

Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with gumbel-softmax. In *International Conference on Learning Representations*, 2017.

Daniel Keysers, Nathanael Schärli, Nathan Scales, Hylke Buiber, Daniel Furber, Sergii Kasber, et al. Measuring compositional generalization: A comprehensive method on realistic data. In *International Conference on Learning Representations*, 2020.

Hyunseo Kim et al. Slot transformer imagination with causality-aware reinforcement learning. In *AAAI Conference on Artificial Intelligence*, 2026.

Thomas Kipf, Ethan Fetaya, Kuan-Chieh Wang, Max Welling, and Richard Zemel. Neural relational inference for interacting systems. In *International Conference on Machine Learning*, 2018.

Sébastien Lachapelle et al. Variational causal dynamics: Discovering modular world models from interventions. In *NeurIPS Workshop on Causal Machine Learning*, 2022.

Yann LeCun. A path towards autonomous machine intelligence. *Open Review*, 2022.

Zhixuan Li et al. Learning factored interactive object-centric world model. In *Advances in Neural Information Processing Systems*, 2025.

Francesco Locatello, Dirk Weissenborn, Thomas Unterthiner, Aravindh Mahendran, Georg Heigold, Jakob Uszkoreit, Alexei Dosovitskiy, and Thomas Kipf. Object-centric learning with slot attention. In *Advances in Neural Information Processing Systems*, 2020.

Heejeong Nam et al. Causal-jepa: Learning world models through object-level latent interventions. *arXiv preprint arXiv:2602.11389*, 2026.

NVIDIA. Cosmos world foundation model platform for physical AI. *arXiv preprint arXiv:2501.03575*, 2025.

OpenAI. Video generation models as world simulators. *Technical Report*, 2024.

Jonathan Richens and Tom Everitt. Robust agents learn causal world models. In *International Conference on Learning Representations*, 2024.

Bernhard Schölkopf, Francesco Locatello, Stefan Bauer, Nan Rosemary Ke, Nal Kalchbrenner, Anirudh Goyal, and Yoshua Bengio. Toward causal representation learning. *Proceedings of the IEEE*, 109(5):612–634, 2021.

Gautam Singh et al. Dreamweaver: Learning compositional world models from pixels. In *International Conference on Learning Representations*, 2025.

Zizhao Wang et al. Dyn-o: Building structured world models with object-centric representations. In *Advances in Neural Information Processing Systems*, 2025.

Ziyi Wu, Nikita Dvornik, Klaus Greff, Thomas Kipf, and Igor Gilitschenski. Slotformer: Unsupervised visual dynamics simulation with object-centric models. In *International Conference on Learning Representations*, 2023.

Zhongwei Yu et al. Learning causal dynamics models in object-oriented environments. In *International Conference on Machine Learning*, 2024.

Hongxin Zhang et al. Compositional world models for embodied multi-agent cooperation. In *International Conference on Learning Representations*, 2025.

A Implementation details

Architecture. State encoder and decoder are 2-layer MLPs with ReLU activation. Edge MLP and type MLP are 3-layer MLPs. Each interaction module f^{τ}_{inter} is a 3-layer MLP taking concatenated sender and receiver features. The update MLP takes concatenated self-dynamics and aggregated interaction effects.

Training. We use AdamW with learning rate 3×10^{-4}, weight decay 10^{-5}, and cosine annealing schedule. Gradient clipping at norm 1.0. Phase 1 (reconstruction only): 10 epochs. Phase 2 (full training): remaining epochs. Autoregressive rollout of 8 steps during training.

Synthetic environment. Objects are colored circles with radius $\in \{4, 6, 8\}$ pixels on a 64×64 canvas. Velocities sampled uniformly from $[-3, 3]$ pixels/frame. Elastic collisions with equal mass. Wall bouncing with perfect reflection.

B Qualitative results

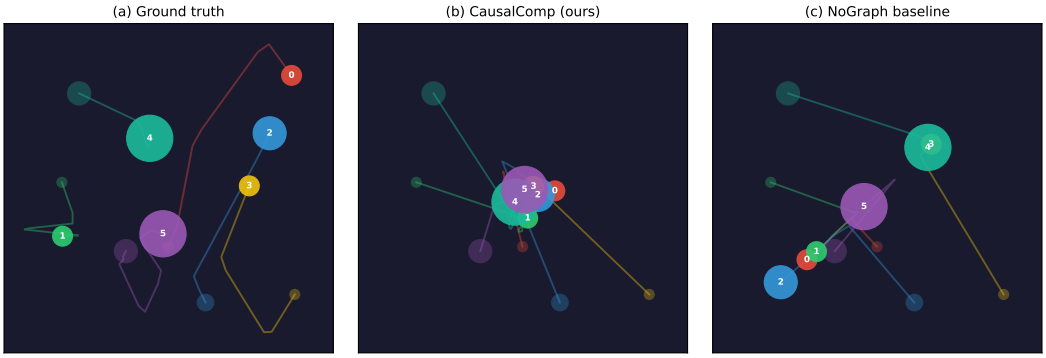

Figure 5: **Trajectory prediction comparison.** (a) Ground-truth trajectories of 6 objects over 16 frames. (b) CAUSALCOMP predictions closely track ground truth, including after collisions. (c) NoGraph baseline shows drift due to inability to model inter-object interactions. Filled circles show final positions; trails show full trajectories.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist"**,
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers**.

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: **[TODO]**

   Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.
   - It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
   - For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
   - If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: **[TODO]**

   Justification: **[TODO]**

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: **[TODO]**

    Justification: **[TODO]**

    Guidelines:

    - The answer [N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
    - If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
    - For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
    - If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: **[TODO]**

    Justification: **[TODO]**

    Guidelines:

    - The answer [N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: **[TODO]**

    Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.